

Problem

Millions manage food allergies, diabetes, celiac disease, hypertension, and other dietary constraints. Grocery platforms know what customers buy — not what they need to avoid. Food discovery remains difficult, time-consuming, and potentially unsafe.


Vision

An AI-powered Food Intelligence platform helping people discover foods aligned with their dietary restrictions, preferences, and health conditions.

Key Value Proposition

The right food, for the right person — powered by intelligence at every layer.
[Github](#)

Solution — Three-Layer Intelligence Architecture



MVP - Current Focus

Food Intelligence — USDA Branded Foods (Cookies Category)

Top 7 Allergens Identified

Milk	Eggs	Tree Nuts	Peanuts	Wheat	Soy	Sesame
------	------	-----------	---------	-------	-----	--------

Success Metrics:

Precision Recall F1 Score Scalability Explainability

Approaches Under Evaluation

- Rule-Based Food Intelligence
- Machine Learning Food Intelligence
- Hybrid (Rules + ML)

Status: dataset acquired · labeller in design · MVP evaluation next

Long-Term Vision

⚠ Food Allergies

💧 Diabetes

🌱 Vegan / Vegetarian Diets

🌾 Celiac Disease

♥ Hypertension

+ Additional Dietary Constraints

Problem

Medications hide top-9 food allergens as inactive ingredients — lactose and starch as binders, soy and egg across drug classes. Patients' documented allergies are never cross-checked at prescribing, so reactions get labeled “idiopathic.”

Vision


A Clinical Decision Support System that sees the problem, stops unsafe prescriptions, and offers a safe, insurance-supported alternative — in time.

Key Insight

The data was always there. The system simply failed to connect it.

[GitHub](#)

Solution — Three-Part Clinical Architecture



MVP - Current Focus

Divya Drishti — Allergen Prevalence in Oral Medications (openFDA)

Initial findings (own analysis)

Dairy 44.86%

Starch 36.58%

Gluten 52.6%*

Soy 20.6%*

n = 42,052 oral medications · * research sample n = 108

Data Sources

- openFDA drug labels
- DailyMed (FDA SPLs)
- RxNorm identifiers

Success Metrics:

Prevalence %

Top Culprit (Mode)

Mfr Distribution

Monthly Freshness

Status: research & problem formulation complete (SIADS 501) · openFDA data acquired · M1 analysis in design

Long-Term Vision

Lakshmana Rekha — EHR / pharmacy alerts (Epic, Cerner)

Celiac — gluten hidden in medications

Policy — evidence base for the ADINA Act

Sanjeevini — formulary-aware safe alternatives

Alpha-gal — gelatin capsules & mammalian excipients

Shared labelling engine with Akshaya Patra

Problem

Finding a name's meaning and origin today means manual Googling, multiple LLM prompts, and summarizing scattered answers yourself — slow, repetitive, and the queries leave your machine.

Vision


A local AI agent (Ollama): type one name, get origin, meaning, and similar names as a clean structured answer — 100% offline, built with enterprise-grade engineering discipline.

Privacy by Design

Names never leave the machine. Zero cloud costs. Zero data leakage.

[Github](#)

Architecture — Guardrails • Agent • Evals



MVP - Current Focus

Name Intelligence Agent — Local Ollama, Guardrailed I/O

Guardrails (pre-LLM)

Profanity block

A-Z only

Non-empty input

Cultural accuracy

Target Metrics:

Schema 100%

Guardrail Block 100%

Hallucination < 5%

Sub-second Response

Evaluation Tiers

• Unit (mocked LLM, deterministic)

• Tester-Agent (agent-on-agent)

• Red-team + performance

Status: PRD & spec complete • repo scaffolded • core build in progress • eval framework designed

Future Extensions

🌐

Multi-language origins & meanings

🗄️

Local SQLite caching — instant repeat lookups

📄

Template for privacy-first local agents

🗣️

Voice — speech-to-text in, audio out

👤

Tester-Agent pattern reused across future agents

📊

Published eval results with every release

Problem

Job platforms match keywords against historical resume data. Career transitions, transferable skills, active learning, and aspirations are invisible — so candidates in transition keep being shown the job they are leaving, not the one they are becoming qualified for.


Vision

An AI career intelligence agent that matches opportunities to a candidate's full trajectory — past experience, present activity, and future intent — and explains every recommendation.

The Thesis

Match jobs to who you are becoming — not just who you have been.
 [Github](#)

Candidate DNA — Three Signal Types, Weighted



MVP - Current Focus

Drishti v1 — Trajectory-Aware Matching (Local, Thin Slice)

v1 components

Trajectory Profile

Curated local JDs

3-signal scoring

Explanations + gaps

Local-First Stack

- Streamlit UI
- Ollama (Qwen) + embeddings
- SQLite + ChromaDB

Target Metrics:

70%+ Relevant

30%+ Non-Obvious Matches

<1% Hallucination

<10s Response

Status: discovery complete · PRD v1 written · Phase A (foundation) in progress

Long-Term Vision

🔑

Live job APIs — RemoteOK, Remotive first

🌐

Multi-agent orchestration (Phase 3)

📊

Skill gap analysis & learning recommendations

🌍

Personalized career roadmaps

📁

GitHub & portfolio signal enrichment

★

“What do I need to do to get there?”

Problem

We rank everything — shopping, doctors, professors — but not the most important choice: where we spend our time, energy, and hope. Applications disappear into a black hole: ghosting, unresponsive processes, zero transparency, no data on how employers actually behave during hiring.


Vision

A standardized, reliable score for how a company treats candidates — visible before a job seeker invests a single hour applying.

The Product Lesson

A real problem + a technical solution ≠ a viable product.
[Github](#)

Investigation — Discovery to Verdict



STATUS — PARKED, NOT ABANDONED

The idea breaks on data — not on AI

Why it fails today

- No journey data
- Biased self-reports
- Seekers won't pay
- Employers won't share
- ATS monetizes employers

Revisit when: **LinkedIn or some other company publishes data** **Regulation mandates transparency** **Employer-funded badge**

Status: investigation complete · published on GitHub · parked with revisit conditions

[Join the LinkedIn discussion →](#)

PM Lessons Learned

- ✓ A real problem ≠ a viable product

✓ Incentives outweigh technology

✓ AI was never the bottleneck
- ✓ Validate data before the AI model

✓ Knowing when NOT to build is a skill

✓ Park with revisit conditions